

VIETNAM NATIONAL UNIVERSITY, HO CHI MINH CITY  
UNIVERSITY OF SCIENCE  
FACULTY OF INFORMATION TECHNOLOGY



## Parallels Processing

---

### Final Report

### Team Contribution

---

Instructor: Phạm Trọng Nghĩa  
Lê Nhựt Nam

Group members:

No.	Full name	Student ID
1	Huỳnh Lê Hải Dương	22127081
2	Nguyễn Triều Khoáng	22127204

## Team planning

### Work contribution

No.	Task name	Student	Deadline	Details	Notes
1	Literature Review & Project Setup	Hải Dương, Triệu Khoáng	01/12/2025	Research autoencoder architecture, CUDA programming. Setup environment, download CIFAR-10 dataset.	Finished (100%)
2	Data Loader Implementation	Hải Dương, Triệu Khoáng	01/12/2025	Implement CIFAR-10 binary file loader, batch generator with shuffling, data normalization.	Finished (100%)
3	CPU Baseline Implementation	Hải Dương	04/12/2025	Implement CPU layers (Conv2D, MaxPool2D, ReLU, Upsampling), CPU autoencoder, training loop with backpropagation.	Finished (100%)
4	GPU Basic Implementation	Hải Dương	07/12/2025	Port layers to GPU, implement CUDA kernels for convolution, pooling, ReLU. GPU autoencoder with memory management.	Finished (100%)
5	GPU Optimization V1 (Shared Memory)	Triệu Khoáng	10/12/2025	Implement shared memory tiling for convolution, reduce global memory accesses, optimize memory access patterns.	Finished (100%)
6	GPU Optimization V2 (Fusion + Streams)	Triệu Khoáng	10/12/2025	Kernel fusion (Conv2D + ReLU + Bias), CUDA streams for computation/data transfer overlap, vectorized memory access.	Finished (100%)
7	Feature Extraction	Triệu Khoáng	13/12/2025	Implement encoder feature extraction, save features to binary format, support both CPU and GPU paths.	Finished (100%)
8	SVM Training Pipeline	Triệu Khoáng	15/12/2025	Integrate cuML SVM, load features and labels, train and evaluate SVM classifier, performance benchmarking.	Finished (100%)
9	Demo Interface	Hải Dương, Triệu Khoáng	16/12/2025	Gradio web interface, load trained models (autoencoder + SVM), real-time image classification, visualization.	Finished (100%)

No.	Task name	Student	Deadline	Details	Notes
10	Performance Analysis & Documentation	Hải Dương, Triệu Khoáng	17/12/2025	Performance profiling and benchmarking, create visualization charts, write Jupyter notebook report, code documentation.	Finished (100%)
11	Video demo	Hải Dương, Triệu Khoáng	18/12/2025	Record and edit video demonstration of the complete system, showcase all features and performance results.	Finished (100%)

## Timeline

No.	Task name	Start date	End date	Duration
1	Literature Review & Project Setup	29/11/2025	01/12/2025	3 days
2	Data Loader Implementation	29/11/2025	01/12/2025	3 days
3	CPU Baseline Implementation	02/12/2025	04/12/2025	3 days
4	GPU Basic Implementation	05/12/2025	07/12/2025	3 days
5	GPU Optimization V1 (Shared Memory)	08/12/2025	10/12/2025	3 days
6	GPU Optimization V2 (Fusion + Streams)	08/12/2025	10/12/2025	3 days
7	Feature Extraction	11/12/2025	13/12/2025	3 days
8	SVM Training Pipeline	13/12/2025	15/12/2025	3 days
9	Demo Interface	14/12/2025	16/12/2025	3 days
10	Performance Analysis & Documentation	29/11/2025	17/12/2025	19 days
11	Video demo	16/12/2025	18/12/2025	3 days

Note: Task 10 (Documentation) is conducted continuously throughout the project. Tasks 8 and 9 can overlap partially.

## Work Distribution Summary

Student	Tasks assigned	Total tasks	Main responsibilities
Huỳnh Lê Hải Dương	1, 2, 3, 4, 9, 10, 11	8	CPU/GPU implementation, optimizations, demo interface, documentation, video demo
Nguyễn Triều Khoáng	1, 2, 5, 6, 7, 8, 9, 10, 11	9	Literature review, data loader, GPU optimization V2, feature extraction, SVM training, demo interface, documentation, video demo

## Key Achievements

- Performance Optimization:** Achieved  $10.91\times$  speedup over CPU baseline through GPU parallelization and optimizations
- Training Time:** Reduced autoencoder training time from 5.7 hours (CPU) to 31.5 minutes (GPU optimized)
- Feature Extraction:** Successfully extracted 8,192-dimensional features from 60,000 CIFAR-10 images in under 23 seconds
- Classification Accuracy:** Achieved 67.53% test accuracy using SVM classifier on extracted features
- Code Quality:** Implemented modular, well-documented codebase with separate CPU and GPU implementations
- Complete Pipeline:** Built end-to-end system from data loading to real-time inference via Gradio interface

## Technologies and Tools Used

- Programming Languages:** C++17, CUDA, Python 3
- GPU Computing:** NVIDIA CUDA Toolkit, cuML (RAPIDS)
- Deep Learning:** Custom autoencoder implementation (no frameworks)
- Machine Learning:** LIBSVM, cuML SVM for classification
- Data Processing:** NumPy, custom CIFAR-10 binary loader
- Visualization:** Matplotlib, Seaborn, Jupyter Notebook
- Web Interface:** Gradio for interactive demo
- Build System:** Makefile for compilation and project management
- Version Control:** Git for collaborative development

## Challenges and Solutions

- **Challenge 1: Upsampling Backward Pass Gradient Accumulation**

- *Problem:* Upsampling forward pass maps one input pixel to four output pixels ( $2\times 2$  upsampling). During backward pass, four gradient values must be accumulated into a single input gradient position, requiring careful indexing and accumulation logic.
- *Solution:* Implemented efficient gradient accumulation by mapping each output gradient back to the corresponding input position using integer division ( $ih = oh \gg 1$ ,  $iw = ow \gg 1$ ) and accumulating gradients correctly without atomic operations.
- *Lesson:* When implementing custom layers, carefully track how forward and backward passes map between tensor dimensions. Visualize the data flow at each step to avoid gradient accumulation errors.

- **Challenge 2: Shared Memory Overhead for Small Images**

- *Problem:* Initial attempt to use shared memory tiling for convolution showed that for small images ( $32\times 32$ ), the overhead of loading tiles and synchronization exceeded the benefit of data reuse.
- *Solution:* Switched to register blocking and weight caching strategy. Used `__restrict__` pointers and loop unrolling to maximize register usage and compiler optimizations. For small kernels ( $3\times 3$ ), this approach proved more efficient than shared memory tiling.
- *Lesson:* Not all optimization techniques work for all problem sizes. Profile and measure - what works for large images may not work for small ones. The optimal strategy depends on image size, kernel size, and available shared memory.

- **Challenge 3: Memory Layout and Coalescing**

- *Problem:* Initial naive implementation had poor memory access patterns, with threads accessing non-consecutive memory locations, leading to uncoalesced memory accesses and poor bandwidth utilization (potentially  $10\times$  reduction in effective bandwidth).
- *Solution:* Reorganized data access patterns to ensure threads in a warp access consecutive memory locations. Used proper indexing schemes and vectorized loads (`float4`) where possible to maximize memory bandwidth. Ensured memory alignment for vectorized operations.
- *Lesson:* Memory access patterns are often more important than raw computation. Uncoalesced accesses can reduce effective bandwidth dramatically. Always design kernels with memory coalescing in mind from the start.

## Project Statistics

Metric	Value
Total lines of code (C++/CUDA)	5,488
Total lines of code (Python)	680
Number of CUDA kernels	15+
Number of CPU functions	20+
Model parameters	751,875
Feature dimension	8,192
Training images processed	50,000
Test images processed	10,000
GPU memory usage (batch size 64)	814 MB